

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

DataSet Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/bareyl/dedicated-bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

```
In [47]: df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

```
In [48]: df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.8
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Bori	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

```
In [49]: df.shape
Out[49]: (10000, 14)
```

```
In [50]: df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+04	650.528800	38.921800	5.012800	76485.689288	1.530200	0.70550	0.515100	100090.239881
std	2886.89568	7.193619e+04	96.653299	10.487808	2.892174	62397.405202	0.581654	0.45584	0.496797	57510.582818
min	1.00000	1.565670e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	1.000000	1.000000	0.00000	0.000000	51002.110000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	7500.25000	1.575232e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	10000.00000	1.581596e+07	850.000000	92.000000	10.000000	250890.090000	4.000000	1.00000	1.000000	199992.480000

```
In [51]: df.isnull()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	False	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns

```
In [52]: df.isnull().sum()
Out[52]: RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
In [53]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   RowNumber            10000 non-null  int64
1   CustomerId           10000 non-null  int64
2   Surname              10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [54]: df.dtypes
Out[54]: RowNumber      int64
CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age             int64
Tenure          int64
Balance         float64
NumOfProducts   int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
Exited          int64
dtype: object
```

```
In [55]: df.columns
Out[55]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'], dtype='object')
```

```
In [56]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis=1) #Dropping the unnecessary columns
```

```
In [57]: df.head()
```

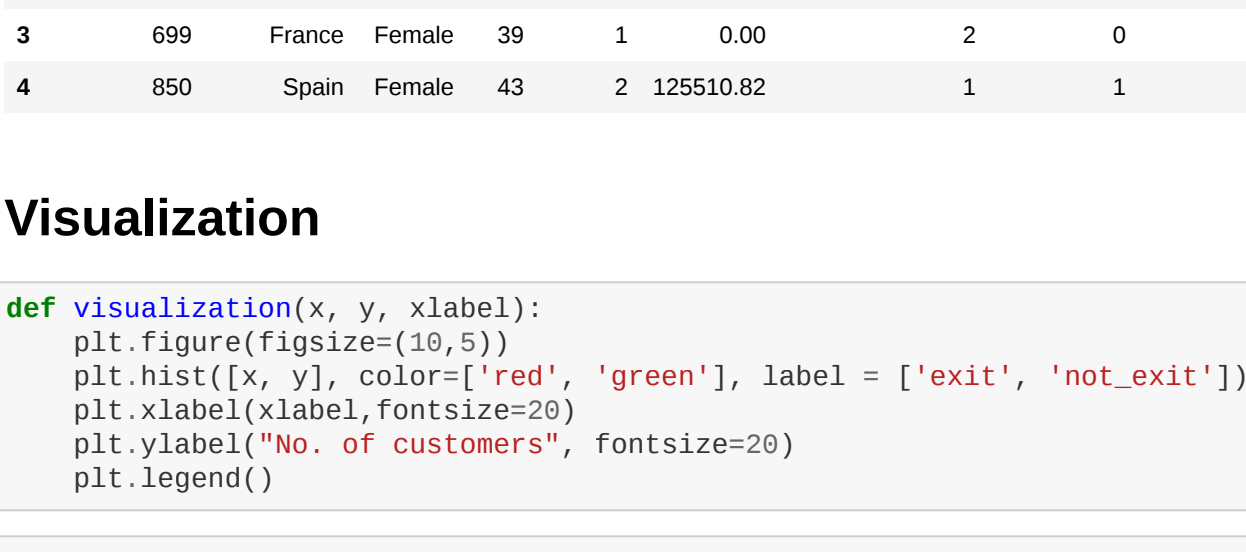
	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

Visualization

```
In [101]: def visualization(x, y, xlabel):
plt.figure(figsize=(10,5))
plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
plt.xlabel(xlabel,fontsize=20)
plt.ylabel("No. of customers", fontsize=20)
plt.legend()
```

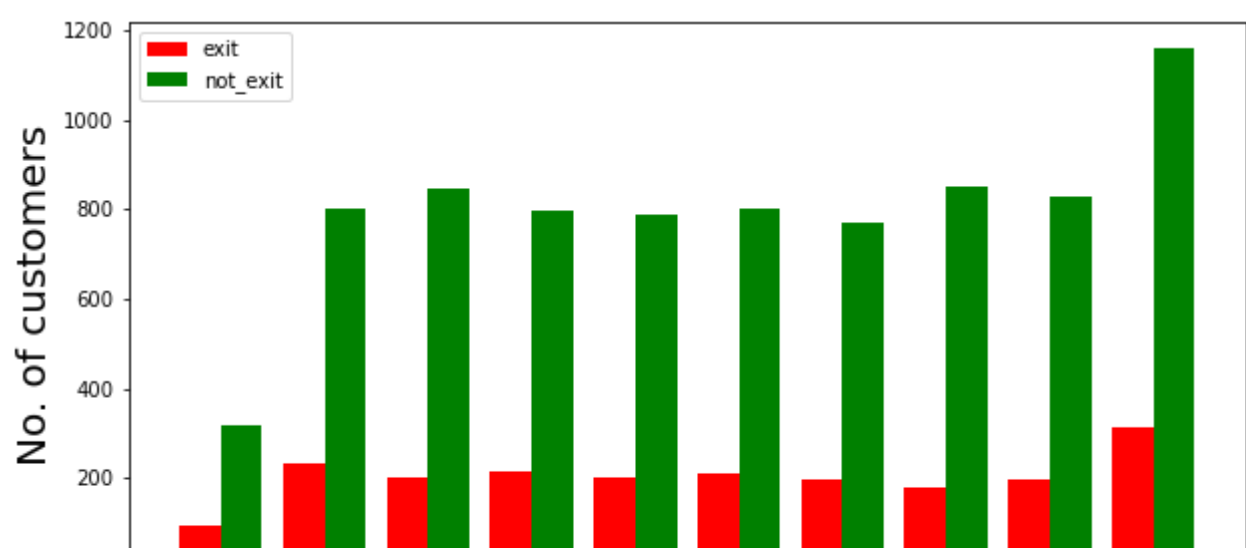
```
In [102]: df_churn_exited2 = df[df['Exited']==1]['Tenure']
df_churn_not_exited2 = df[df['Exited']==0]['Tenure']
```

```
In [103]: visualization(df_churn_exited2, df_churn_not_exited2, "Tenure")
```



```
In [105]: df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [106]: visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

```
In [59]: X = df[['CreditScore','Gender','Age','Tenure','Balance','NumOfProducts','HasCrCard','IsActiveMember','EstimatedSalary']]
states = pd.get_dummies(df[['Geography']],drop_first = True)
gender = pd.get_dummies(df[['Gender']],drop_first = True)
```

```
In [61]: df = pd.concat([df,gender,states], axis = 1)
```

Splitting the training and testing Dataset

```
In [62]: df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Male	Germany	Spain
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1	0	0	0
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0	0	0	1
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1	0	0	0
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0	0	0	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0	0	0	1

```
In [63]: X = df[['CreditScore','Age','Tenure','Balance','NumOfProducts','HasCrCard','IsActiveMember','EstimatedSalary'],'Male','Germany','Spain']
```

```
In [64]: y = df['Exited']
```

```
In [65]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
In [66]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [67]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [68]: X_train
```

```
Out[68]: array([[ 4.56838557e-01, -0.45594735e-01, 1.58341939e-03, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01],
 [-2.07591864e-02, -2.77416637e-01, 3.47956411e-01, ...,
 -1.09507222e+00, -5.81969145e-01, -1.74334114e+00],
 [-1.66115021e-01, 1.82257167e+00, -1.38390855e+00, ...,
 -1.09507222e+00, -5.81969145e-01, -5.73611209e-01],
 ...,
 [-3.63383654e-01, -4.68324665e-01, 1.73344838e+00, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01],
 [ 4.6722117e-01, -1.42264640e+00, -1.38707539e+00, ...,
 9.13181783e-01, -5.81969145e-01, -1.74334114e+00],
 [-8.82511636e-01, 2.95307447e-01, -6.91162564e-01, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01]])
```

```
In [69]: X_test
```

```
Out[69]: array([[ 3.63395520e-01, 1.99853433e-01, 1.58341939e-03, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01],
 [-4.15243057e-02, 4.86215475e-01, 1.58341939e-03, ...,
 -1.09507222e+00, -5.81969145e-01, -1.74334114e+00],
 [-1.87823736e+00, -3.72870651e-01, -1.38390855e+00, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01],
 ...,
 [-0.82182526e-01, -5.63778679e-01, -1.73028154e+00, ...,
 -1.09507222e+00, -5.81969145e-01, -5.73611209e-01],
 [ 1.15189594e+00, -6.59232693e-01, 1.73344838e+00, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01],
 [-5.19122049e-01, 1.84399419e-01, 1.73344838e+00, ...,
 9.13181783e-01, -5.81969145e-01, -5.73611209e-01]])
```

Building the Classifier Model using Keras

```
In [70]: import keras #keras is the wrapper on the top of tensorflow
#Can use Tensorflow as well but won't be able to understand the errors initially.
```

```
In [71]: from keras.models import Sequential #To create sequential neural network
from keras.layers import Dense #To create hidden layers
```

```
In [72]: classifier = Sequential()
```

```
In [74]: #To add the layers
#Dense helps to connect the neurons
#Input Dimension means we have 11 features
# Units is to create the hidden layers
#Uniform helps to distribute the weight uniformly
classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "uniform"))
```

```
In [75]: classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) #Adding second hidden layers
```

```
In [76]: classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) #Final neuron will be having sigmoid function
```

```
In [77]: classifier.compile(optimizer='adam',loss = 'binary_crossentropy',metrics = ['accuracy']) #To compile the Artificial Neural Network. Used Binary Crossentropy as we just have only two output
```

```
In [79]: classifier.summary() #3 layers created. 6 neurons in 1st,8neurons in 2nd layer and 1 neuron in last
```

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 6)	72
dense_4 (Dense)	(None, 6)	42
dense_5 (Dense)	(None, 1)	7
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		

```
In [80]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to training dataset
```

```
Epoch 1/50
700/700 [=====] - 0s 674us/step - loss: 0.4293 - accuracy: 0.7947
Epoch 2/50
700/700 [=====] - 0s 647us/step - loss: 0.4239 - accuracy: 0.7947
Epoch 3/50
700/700 [=====] - 0s 657us/step - loss: 0.4203 - accuracy: 0.8067
Epoch 4/50
700/700 [=====] - 0s 664us/step - loss: 0.4167 - accuracy: 0.8260
Epoch 5/50
700/700 [=====] - 0s 674us/step - loss: 0.4153 - accuracy: 0.8287
Epoch 6/50
700/700 [=====] - 0s 653us/step - loss: 0.4137 - accuracy: 0.8318
Epoch 7/50
700/700 [=====] - 0s 658us/step - loss: 0.4125 - accuracy: 0.8317
Epoch 8/50
700/700 [=====] - 1s 842us/step - loss: 0.4116 - accuracy: 0.8306
Epoch 9/50
700/700 [=====] - 0s 671us/step - loss: 0.4180 - accuracy: 0.8331
Epoch 10/50
700/700 [=====] - 0s 682us/step - loss: 0.4163 - accuracy: 0.8326
Epoch 11/50
700/700 [=====] - 0s 690us/step - loss: 0.4093 - accuracy: 0.8337
Epoch 12/50
700/700 [=====] - 0s 688us/step - loss: 0.4087 - accuracy: 0.8339
Epoch 13/50
700/700 [=====] - 0s 675us/step - loss: 0.4081 - accuracy: 0.8341
Epoch 14/50
700/700 [=====] - 1s 722us/step - loss: 0.4071 - accuracy: 0.8331
Epoch 15/50
700/700 [=====] - 0s 811us/step - loss: 0.4065 - accuracy: 0.8341
Epoch 16/50
700/700 [=====] - 1s 771us/step - loss: 0.3984 - accuracy: 0.8367
Epoch 17/50
700/700 [=====] - 0s 695us/step - loss: 0.3976 - accuracy: 0.8374
Epoch 18/50
700/700 [=====] - 0s 669us/step - loss: 0.3972 - accuracy: 0.8373
Epoch 19/50
700/700 [=====] - 0s 670us/step - loss: 0.3970 - accuracy: 0.8376
Epoch 20/50
700/700 [=====] - 1s 720us/step - loss: 0.3972 - accuracy: 0.8376
Epoch 21/50
700/700 [=====] - 0s 675us/step - loss: 0.3965 - accuracy: 0.8367
Epoch 22/50
700/700 [=====] - 0s 680us/step - loss: 0.3961 - accuracy: 0.8364
Epoch 23/50
700/700 [=====] - 0s 685us/step - loss: 0.3962 - accuracy: 0.8379
Epoch 24/50
700/700 [=====] - 1s 771us/step - loss: 0.3960 - accuracy: 0.8370
Epoch 25/50
700/700 [=====] - 1s 1ms/step - loss: 0.3963 - accuracy: 0.8366
Epoch 26/50
700/700 [=====] - 1s 746us/step - loss: 0.3962 - accuracy: 0.8373
Epoch 27/50
700/700 [=====] - 1s 823us/step - loss: 0.3950 - accuracy: 0.8384
Epoch 28/50
700/700 [=====] - 1s 759us/step - loss: 0.3956 - accuracy: 0.8361
Epoch 29/50
700/700 [=====] - 0s 773us/step - loss: 0.3949 - accuracy: 0.8369
Epoch 30/50
700/700 [=====] - 0s 695us/step - loss: 0.3953 - accuracy: 0.8366
Epoch 31/50
700/700 [=====] - 0s 701us/step - loss: 0.3952 - accuracy: 0.8369
Epoch 32/50
700/700 [=====] - 0s 707us/step - loss: 0.3952 - accuracy: 0.8366
Epoch 33/50
700/700 [=====] - 0s 680us/step - loss: 0.3955 - accuracy: 0.8376
Epoch 34/50
700/700 [=====] - 0s 665us/step - loss: 0.3957 - accuracy: 0.8373
Epoch 35/50
700/700 [=====] - 0s 680us/step - loss: 0.3947 - accuracy: 0.8371
Epoch 36/50
700/700 [=====] - 0s 681us/step - loss: 0.3944 - accuracy: 0.8371
Epoch 37/50
700/700 [=====] - 0s 678us/step - loss: 0.3947 - accuracy: 0.8383
Epoch 38/50
700/700 [=====] - 1s 869us/step - loss: 0.3944 - accuracy: 0.8370
```

```
Out[89]: <tensorflow.python.keras.callbacks.History at 0x1f1eb93df0>
```

```
In [90]: y_pred=classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result
```

```
In [97]: from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
```

```
In [92]: cm = confusion_matrix(y_test,y_pred)
```

```
In [93]: cm
```

```
Out[93]: array([[2328, 72],
 [ 425, 175]], dtype=int64)
```

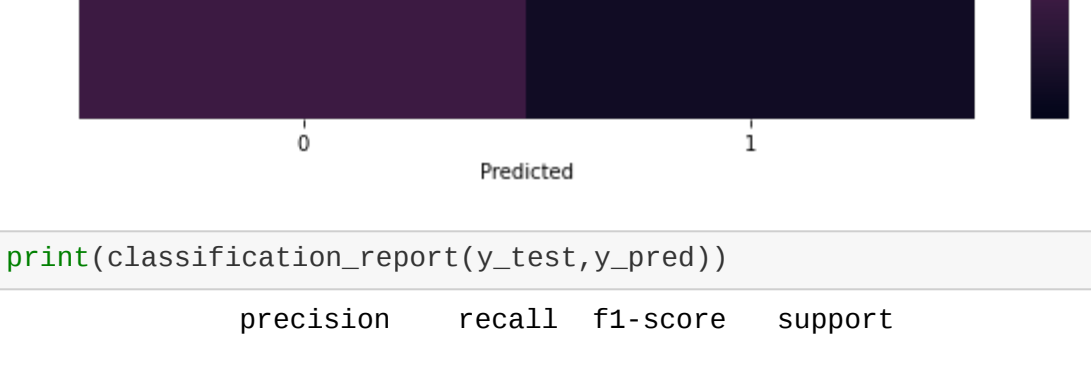
```
In [94]: accuracy = accuracy_score(y_test,y_pred)
```

```
In [95]: accuracy
```

```
Out[96]: 0.8343333333333334
```

```
In [98]: plt.figure(figsize = (10,7))
sns.heatmap(cm,annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

```
Out[98]: Text(69.0, 0.5, 'Truth')
```



```
In [100]: print(classification_report(y_test,y_pred))
```

accuracy			0.83	3000
macro avg	0.78	0.63	0.66	3000
weighted avg	0.82	0.83	0.81	3000